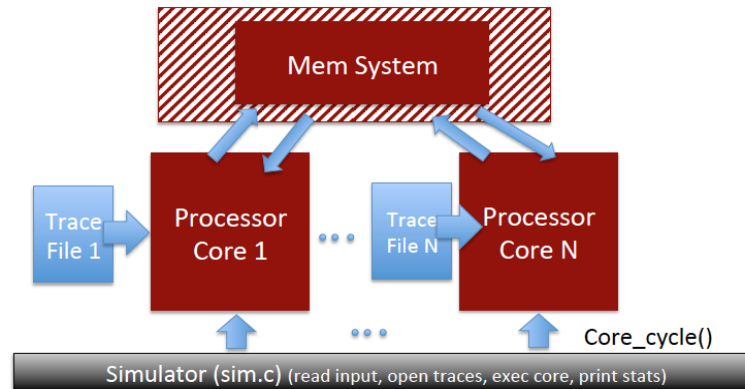


CS 4290 / CS 6290 / ECE 4100 / ECE 6100
Advanced Computer Architecture

Lab 4: CMP Memory System Design (10 (+ 5) Pts)

Parts A + B + C Due: Tuesday, November 22, 2022 (11:55 pm)

Part D (+ E + F) Due: Thursday, December 1, 2022 (11:55 pm)



This is an individual assignment. You can discuss this assignment with other classmates but you should code your assignment individually. *You are **NOT** allowed to see the code of (or show your code to) other students or of past submissions that may be available online.* We will be using software that detects code similarity. If you think you collaborated closely enough with another student that could lead to more code similarities than would normally be expected, please declare it upfront, at time of submission (as a comment in your Canvas submission). *Please note that such declaration is not granting you an excuse for copying code.*

You are given reference execution times for every part. There is no execution time limit; you won't be graded on how long your code takes to run. If your code takes way more time (e.g. 5x) to execute than the reference time, you may be doing something very inefficiently and may want to reconsider your implementation. You can time your experiments using the [time command](#).

OBJECTIVE

The objective of the fourth (and last) programming assignment is to build and evaluate a multi-level cache simulator with DRAM-based main memory. The system will then be extended to incorporate multiple cores, where each core has a private L1 (I and D) cache, and a shared L2 cache. Misses and writebacks from the shared L2 cache are

served by a DRAM-based main memory consisting of 16 banks and per-bank row buffers.

PROBLEM DESCRIPTION

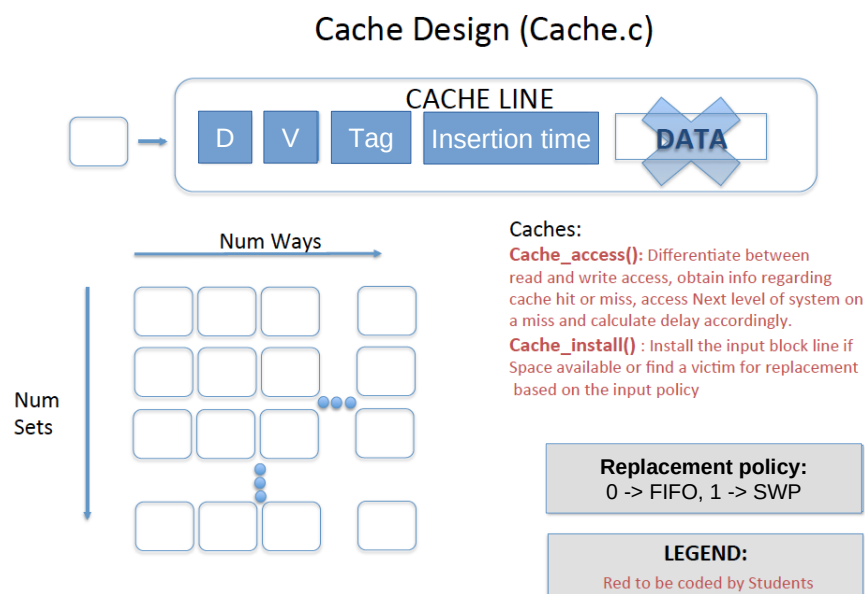
The figure above shows the multicore system with the memory hierarchy that you will build as part of this lab. We will build the simulator for this system in two phases. The first phase (A, B, C) is for a single-core system and the second phase (D, E, F) extends the system to support multiple cores.

Tip: The complexity of the system you must implement in this lab is lower than Lab 3. However, you are given a more loosely defined code structure. Hence, you should plan to invest more time upfront, designing your code's skeleton.

Part A: Design a Standalone Cache Module (3 points)

Reference execution time: 9-10 sec / trace

In this part, you will build a cache model and estimate the cache miss ratio. You need to implement your own data structure for the cache (name it **Cache**) in cache.{cpp, h} and leave all other files untouched (except for run.sh). The cache functions must be written to work for different associativity, line size, cache size, and replacement policies. Refer to Appendix A for more details on the cache implementation.



As we are only interested in cache hit/miss information we will not be storing data values in our cache. We provide you with traces of 100M instructions from three SPEC2006 benchmarks: bzip2, lbm, and libq.

Each trace record consists of 4 bytes of instruction PC, 1 byte of instruction type (0: ALU, 1:LOAD, 2:STORE) and 4 bytes of virtual address for LD/ST instructions. We are providing you with the trace reader and the ability to change cache parameters from the command line. The trace reader calls the memory system for LD and ST instructions and the function in the memory system in turn calls the **cache_access** function for the DCACHE. If the line is not found in the DCACHE, the memsys function calls the **cache_install** function for the DCACHE, which in turn calls **cache_find_victim**. For this Lab, you should implement a FIFO replacement policy. In the end, the memory system also calls the **cache_print_stats** function for the DCACHE. The cache print stats functions are already written for you.

Your objective is to implement three functions in **cache.cpp**:

1. **cache_access** – if the line is present in the cache, return HIT.
2. **cache_install** – install the line in the cache and track the evicted line.
3. **cache_find_victim** – find the victim to be evicted.

Refer to Appendix A for more details on the Cache implementation.

Part B: Multi-level Cache (2 points)

Reference execution time: 12-13 sec / trace

You will implement an instruction cache (ICACHE) and a data cache (DCACHE), and connect them to a unified L2 cache and DRAM. You will also estimate timing for each request in this part.

Your objective is to write two functions in **memsys.cpp**:

1. **memsys_access_modeBC**, which returns the delay required to service a given request.
2. **memsys_L2_access**, which is called by **memsys_access_modeBC** and returns the delay for servicing a request that accesses the L2 cache.

Note that for the purpose of calculating delay, you can assume that writebacks are completed off the critical path and hence do not account for the accumulation of delay for a read request.

NOTE: All caches are Write-back and Allocate-On-Miss. For this part assume a fixed DRAM Latency of 100 cycles.

Part C: Implementing a Simple DRAM (3 points)

Reference execution time: 12-13 sec / trace

For Part C, you will model a simple row buffer for the DRAM memory system. You will assume that the DRAM memory has 16 banks and the memory system could follow either open page or close page policy. Address mapping should use cache block interleaving: consecutive cache blocks should be mapped to consecutive banks.

Your objective is to implement the DRAM class in `dram.{cpp, h}`. `memsys_access_modeCDE` returns the DRAM delay for **both open page and close page policy**. This can be controlled by a command line argument. By default, it should return the fixed value (100 cycles) as mentioned in the previous part.

Refer to Appendix B for more details on the DRAM implementation.

Part D: Making the system Multicore (2 points)

Reference execution time: 26-27 sec / mix

You will assess the effectiveness of your memory system for a multicore processor. In particular, you will implement a two-core processor and evaluate your design for three mix workloads: Mix1 (libq-bzip2), Mix2(libq-lbm), and Mix3(bzip2-lbm). You will model simple FIFO replacement in the shared L2 cache and report the weighted speedup under FIFO replacement. Pay attention to the memory traffic and the memory row buffer hit rate for the mix workload compared to the isolated workloads from Part C.

Note: There is no need for a coherence protocol in this lab. We use a pseudo-translation layer (`memsys_convert_vpn_to_pfn()`) that produces physical addresses unique to each core. Since there is no physical address overlap between cores, there is no need to maintain coherence information.

Your objective is to write two functions in `memsys.cpp`:

1. `memsys_access_modeDE`, which returns the delay required to service a given request.

2. [memsys_L2_access_multicore](#), which is called by `memsys_access_modeDE` and returns the delay for servicing a request that accesses the L2 cache.

Part E: Implement Static Way Partitioning (2 points Extra Credit)

Reference execution time: 26-27 sec / mix

Shared caches can cause an ill-behaved application to completely consume the capacity of the shared cache and cause significant performance degradation to the neighboring application. This can be mitigated with way partitioning, whereby each core can be given a fixed quota of ways per set. In this part you will implement static way partitioning for a system consisting of two cores. We will provide “SWP_core0ways” as the quota for core 0 (the remaining $N - \text{SWP_core0ways}$ becomes the quota of core 1).

Note 1: Invalid lines are still the prime candidates for selection, even if selecting an invalid line as victim results in a core’s quota being exceeded. This means that a core can temporarily exceed its way-partition quota.

Note 2: When selecting from multiple valid lines in a partition, use FIFO as your replacement policy.

Your objective is to update [cache_find_victim](#) in `cache.cpp` to handle SWP.

Part F: Dynamic Way Partitioning (3 points Extra Credit)

Implement any partitioning scheme different from Static Way Partitioning. One option is to implement a [Utility-Based Cache Partitioning Scheme](#). You can read, understand the scheme and implement it in whatever way you wish to. You are free to implement anything else by looking through other technical papers, or even try something of your own.

You are allowed to add any structure you would like to implement the scheme. The only restrictions are that you are not allowed to change the size, associativity, number of sets or block size for any of the caches. All you have to do is to implement your choice of a non-static partitioning scheme, which would decide how many ways of the L2 cache are dedicated to which core. *Note that no code debugging help by the TAs will be provided for this part.*

You will implement this scheme ensuring the code behaves as expected for the other modes of the Lab as well. Your scheme should be invoked by setting the command line parameters of mode as 4 and L2repl as 3, as presented to you in the runall.sh file.

Write a report briefly describing the technique you have implemented, references, your idea on why it would provide better performance and the L2 miss ratios observed by running your code on the given traces.

The grading for this part **will depend on your report**, your understanding of the new scheme and **your implementation**.

+3 points if your implementation is correct and, compared to parts D and E (for core0ways=50%):

- it does not increase the L2 miss ratio for any of the three instruction mixes, and
- it lowers L2 miss ratio by more than 1% for any one of the three instruction mixes.

+1 point if your implementation is correct but you do not beat Parts D and E for any instruction mix.

How to run the simulator:

1. `./sim -h` (to see the simulator options)
2. `./sim mode 1 ../traces/*.mtr.gz` (to test the default configuration)
3. `../scripts/runall.sh` runs all configurations for all traces

WHAT TO SUBMIT FOR Parts A+B+C:

For phase 1 (parts A, B, C) you will submit two compressed folders:

1. src_ABC.tar.gz, containing the src folder with all source code files and the makefile:

src_ABC.tar.gz:

src/*.cpp

src/*.h

src/makefile

2. results_ABC.tar.gz, which is a tarball of your results for the three parts, for each of the three trace files (12 .res files in total).

WHAT TO SUBMIT FOR Part D (+E+F):

For phase 2 (parts D, E, F) you will submit two compressed folders and a pdf report:

1. **src_DEF.tar.gz**, containing the src folder with all source code files and the makefile:

src_DEF.tar.gz:

src/*.cpp

src/*.h

src/makefile

2. **results_DEF.tar.gz** (which is a tarball of your results for the three parts, for each of the three trace files)

3. **DWP_report.pdf** (for Part F)

REFERENCE MACHINE:

As in the previous labs, we will use the virtual machine **oortcloud.cc.gatech.edu**.

Before submitting your code ensure that it compiles on this machine and generates the desired output, without any extra printf statements. Please do not upload logs, object files (*.o) and binaries.

Please follow all the submission instructions.

If you do not follow the submission instructions and file names, you will not receive full credit. If your code does not compile and run correctly on oortcloud, you will not receive credit.

Hints:

1. We have provided `cycle_count` as a means for tracking the time, which can be used to implement the FIFO replacement policy

2. You will need the `last_evicted` line for Part B, when you must schedule a writeback from the Dcache to the L2cache, and from the L2cache to DRAM.

Appendix A: Cache Model implementation

In the “src” directory, you need to update two files:

- `cache.h`
- `cache.cpp`

The following data structures may be needed for completing the Lab. You are free to use any name, any function (with any number of arguments) to implement a cache. You are free to deviate from these structural definitions as well.

1. “**Cache Line**” structure (`Cache_Line`), will have the following fields:
 - **Valid**: denotes if the cache line is indeed present in the Cache
 - **Dirty**: denotes if the latest data value is present only in the local Cache
 - **Tag**: denotes the conventional higher-order address bits beyond Index

- **Core ID:** needed to identify the core to which a cache line (way) is assigned to in a multicore scenario (required for Part D, E, F)
 - **Insertion Time:** to keep track of when each line was inserted, which helps with the FIFO replacement policy
2. **“Cache Set”** structure (Cache_Set), will have:
 - **Cache_Line Struct** (replicated “# of Ways” times, as in an array/list)
 3. The overarching **“Cache”** structure should have:
 - Cache_Set Struct (replicated “#Sets” times, as in a list/array)
 - # of Ways
 - Replacement Policy
 - # of Sets
 - Last evicted Line (Cache_Line type) to be passed on to next higher cache hierarchy for an install if necessary

Status Variables (Mandatory variables required for generating the desired final reports as necessary. Aimed at complementing your understanding of the underlying concepts):

- **stat_read_access:** Number of read (lookup accesses do not count as READ accesses) accesses made to the cache
- **stat_write_access:** Number of write accesses made to the cache
- **stat_read_miss:** Number of READ requests that lead to a MISS at the respective cache
- **stat_write_miss:** Number of WRITE requests that lead to a MISS at the respective cache
- **stat_dirty_evicts:** Count of requests to evict DIRTY lines

Take a look at “types.h” to choose appropriate datatypes.

NOTE: Don’t change the print order/format/names of the cache_print_stats function.

Maximum # of WAYS should be 16.

The variables from sim.cpp might be useful:

SWP_CORE0_WAYS, cycle (You can use this as timestamp for FIFO).

You can access them using 'extern'.

Appendix B: (DRAM model)

In the “src” directory, you need to update two files:

- dram.h
- dram.cpp

The following data structures may be needed for completing the Lab. You are free to use any name, any function (with any number of arguments) to implement the DRAM. You are free to deviate from these structural definitions as well.

1. **“Row Buffer”** Entry structure (Rowbuf_Entry) can have following entries:
 - Valid
 - Row ID (If the entry is valid, which row)
2. **“DRAM”** structure can have the following fields:
 - Array of Rowbuf Entry (Maximum could be 256)

Status Variables (Mandatory metric variables required to be implemented and updated as necessary):

- **stat_read_access**: Number of read (lookup accesses do not count as READ accesses) accesses made to the DRAM
- **stat_write_access**: Number of write accesses made to the DRAM
- **stat_read_delay**: keeps track of cumulative DRAM read latency for subsequent incoming READ requests to DRAM (only the latency spent at DRAM module)
- **stat_write_delay**: keeps track of cumulative DRAM write latency for subsequent incoming WRITE requests to DRAM (only the latency paid at DRAM module)

NOTE: Don't change the print order/format/names of the dram_print_stats function.

DRAM Latencies for this part:

ACT (RAS)	45
CAS	45
PRE	45
BUS	10

Note: A single DRAM access incurs the indicated BUS latency only once (i.e., do not add the bus latency twice, assuming you pay the latency once to get to the DRAM and once for data to return from DRAM).

Row Buffer Size = 1024

DRAM Banks = 16

The following variables from sim.cpp might be useful: SIM_MODE, CACHE_LINESIZE, DRAM_PAGE_POLICY. You can access them using 'extern'.